Q1.
Client-Server Model in Distributed Systems

The **client-server model** is a widely used architectural design in distributed systems where there is a clear division of roles between two types of entities: the **clients** and the **servers**. This model enables multiple clients to request services from servers, which respond with the requested data or service.

Structure of Interaction Between Clients and Servers:

1. **Clients:**

   - **Role:** Clients are the consumers of services. They initiate requests for data, computation, or services from servers. Typically, clients are the end-user devices or applications.
   - **Example:** A web browser (client) requesting a webpage from a web server.
   - **Characteristics:**
     - Clients are generally simpler, with less processing power than servers.
     - They handle the user interface and interact with the end user directly.
     - They do not share resources with other clients.

2. **Servers:**

   - **Role:** Servers are powerful machines or programs designed to provide resources, services, or data to clients. They listen for incoming requests from clients and respond accordingly.
   - **Example:** A web server hosting websites or a database server responding to SQL queries.
   - **Characteristics:**
     - Servers often have greater processing power and storage to manage large workloads.
     - They provide centralized services to multiple clients.
     - They run continuously to be available for client requests at any time.

Interaction Process:

1. **Request-Response Model:**

   - The interaction between clients and servers follows a **request-response cycle**. Clients send requests to the server, which processes the request and returns the appropriate response (e.g., data or a service).

2. **Network Communication:**

   - Communication between clients and servers happens over a network (LAN, WAN, or the internet), using standardized communication protocols such as **HTTP**, **TCP/IP**, or **RPC**.
   - Clients usually interact with servers through APIs or service endpoints that expose the server's functionality.

3. **State Management:**

- **Stateless Servers:** Servers like web servers typically do not retain information about clients between requests. This means each request is processed independently (e.g., HTTP servers).
- **Stateful Servers:** Some servers maintain client-specific information across requests (e.g., a database server handling user sessions).

Advantages of the Client-Server Model:

- **Centralized Management:** Resources such as data, applications, or services are managed centrally on the server, which simplifies administration and updates.
- **Scalability:** More clients can be added without significantly affecting the server. Servers can also be upgraded or distributed to handle larger loads.
- **Security:** Security measures such as authentication and access control are centralized at the server level, making it easier to enforce and manage.

Limitations of the Client-Server Model:

- **Single Point of Failure:** If the server fails, clients cannot access the requested resources.
- **Scalability Limits:** A single server might become a bottleneck under heavy load unless distributed server solutions are implemented.

Example Use Cases:

- **Web Services:** A web browser (client) interacting with a web server to retrieve webpages.
- **Database Systems:** A database client querying a server for data.

(Add TPS , EAI related thory a bit) (Client server diag)

Q2.

Relationship Between Middleware and the Client-Server Model in Distributed Systems

Middleware plays a crucial role in enhancing the **client-server model** by acting as a layer between the client and server to facilitate communication, manage complexity, and provide essential services. In distributed systems, middleware abstracts many of the underlying complexities involved in distributed computing and enables seamless interaction between the client and server.

An illustration of how middleware works in distributed systems
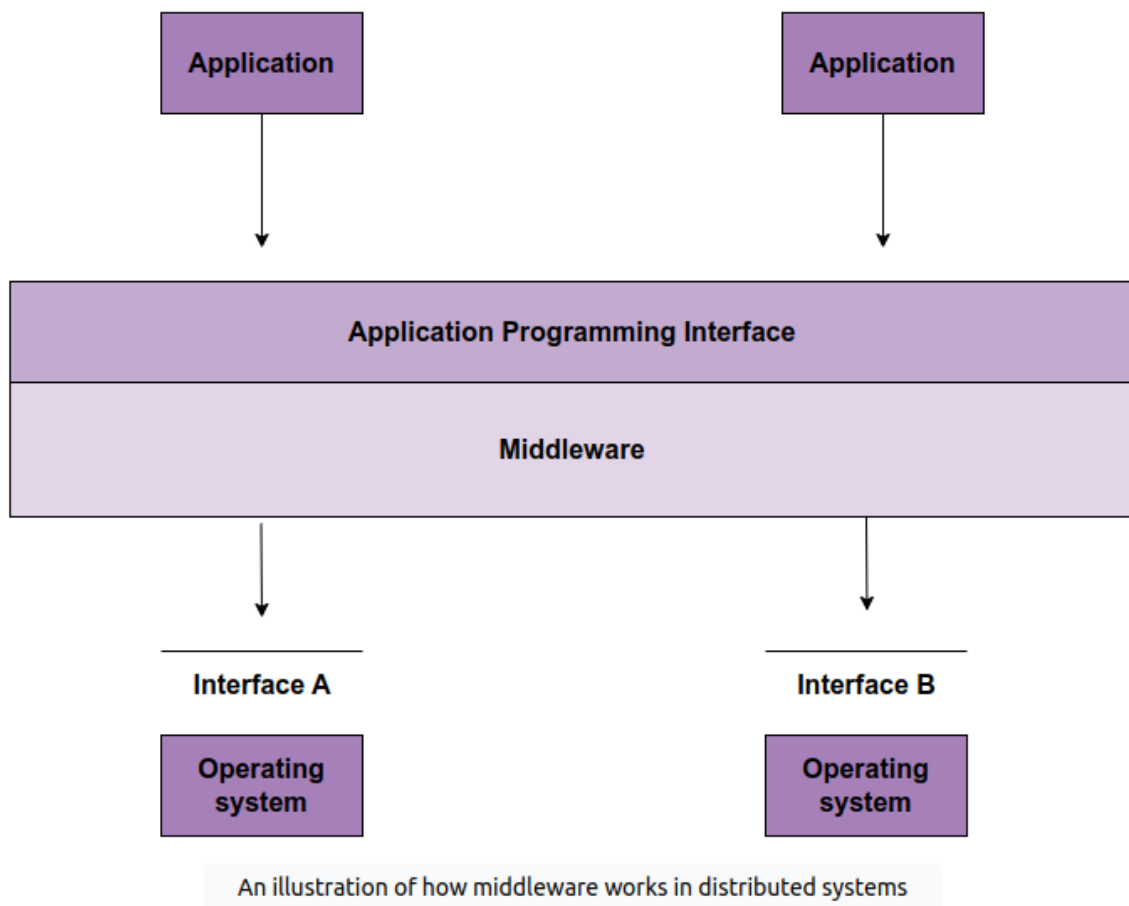
1. Role of Middleware in the Client-Server Model:

Middleware is software that provides a set of services to enable the interaction between clients and servers. It hides the complexity of the underlying network and system architecture, allowing developers to focus on the application logic without worrying about the technical details of the communication process.

Key Functions of Middleware:

- **Communication Management:** Middleware handles the communication between the client and the server, making sure that messages are delivered reliably and efficiently.
- **Abstraction:** It abstracts the underlying hardware and network infrastructure, allowing the client and server to interact without worrying about the specifics of data transmission or protocols.
- **Interoperability:** Middleware facilitates interaction between heterogeneous systems, enabling clients and servers running on different platforms or written in different programming languages to communicate seamlessly.
- **Service Management:** Middleware provides common services like authentication, security, transactions, load balancing, and data access to make the system more robust and scalable.

2. Enhancing the Functionality of the Client-Server Architecture:

Middleware significantly enhances the functionality of the client-server model by simplifying the development process and improving the system's efficiency, security, and scalability.

How Middleware Enhances Client-Server Functionality:

1. **Simplifies Communication:**

   - Middleware provides standardized communication mechanisms (like Remote Procedure Calls [RPC], Message-Oriented Middleware [MOM], or Object Request Brokers [ORB]) that enable clients and servers to communicate easily without needing to know the network specifics.
   - Example: A client does not need to manage low-level network protocols (e.g., TCP/IP); middleware ensures that requests are routed correctly to the server.

2. **Transparency:**

   - Middleware offers **distribution transparency**, meaning that clients are unaware of the physical location of servers or the complexity of the distributed environment. This ensures that distributed systems feel like unified, coherent systems to the user.
   - Example: In a distributed database, middleware ensures that the client can access data from any server, regardless of its actual location.

3. **Security:**

   - Middleware often provides built-in security services like **authentication**, **encryption**, and **authorization**. This ensures that only legitimate clients can access server resources, enhancing the security of the overall system.
   - Example: Middleware can manage secure communication channels (using protocols like SSL/TLS) to protect data transmitted between the client and server.

4. **Scalability and Load Balancing:**

   - Middleware can manage load distribution across multiple servers, ensuring that requests are evenly balanced and preventing any single server from becoming a bottleneck.
   - Example: Middleware can direct client requests to the least busy server in a cluster, improving performance and scalability in high-traffic environments.

5. **Fault Tolerance:**

   - Middleware improves **fault tolerance** by providing mechanisms for handling server failures or network disruptions. It can ensure that client requests are rerouted or retried if a server becomes unavailable.
   - Example: Middleware can automatically redirect client requests to a backup server if the primary server goes down.

6. **Interoperability and Heterogeneity:**

   - Middleware provides a common interface, allowing clients and servers that operate on different platforms, operating systems, or programming languages to work together.
   - Example: A Java-based client can interact with a server written in C++ through middleware that translates the communication between them.

3. Examples of Middleware in Client-Server Systems:

- **CORBA (Common Object Request Broker Architecture):** Enables objects written in different languages to communicate in a distributed environment.
- **Web Services Middleware:** Provides an interface for applications to interact using web-based protocols (e.g., SOAP, REST).
- **Message-Oriented Middleware (MOM):** Queues messages between clients and servers, ensuring reliable delivery even in asynchronous communication.

Q3.

Introduction

In distributed systems, communication between different nodes or devices over a network is critical. Two major network models that define how data is transmitted between these nodes are the **TCP/IP model** and the **OSI model**. Both models provide layered frameworks to support **distributed communication**, but they differ in structure, complexity, and real-world implementation. Comparing them helps us understand their roles in ensuring reliable communication across distributed systems.

---

1. Overview of TCP/IP and OSI Models

TCP/IP Model:

- **Also Known As**: Internet Protocol Suite
- **Layers**: Application, Transport, Internet, Network Access (Link)
- **Purpose**: Designed for practical, real-world communication, forming the backbone of the internet and many distributed systems.

OSI Model:

- **Layers**: Application, Presentation, Session, Transport, Network, Data Link, Physical
- **Purpose**: A conceptual model that standardizes network functions, often used as a reference in understanding network communication.
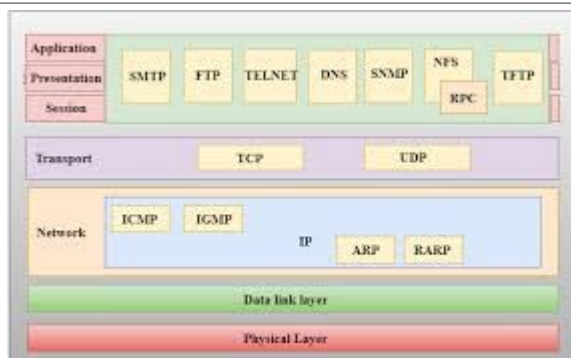
2. Comparison Between the TCP/IP and OSI Models

Aspect	TCP/IP Model	OSI Model
Layering Approach	4 layers, combining certain functions (e.g., Application, Presentation, Session into one Application layer).	7 layers, with distinct roles for each communication function.
Purpose	Practical and implementation-driven. Used in real-world networks, including distributed systems.	A theoretical model for standardizing network functions.
Transport Protocol	TCP (reliable) and UDP (unreliable) for communication.	Provides both reliable and unreliable transport services conceptually.
Flexibility and Protocol Dependency	Protocols like HTTP, FTP are tightly coupled with layers.	More flexible, but less practical, as it separates functions across more layers.
Usage in Distributed Systems	Dominates real-world use in distributed systems (cloud, web services, etc.).	Rarely implemented in full; used as a reference for teaching and protocol design.

3. Support for Distributed Communication

TCP/IP Model in Distributed Systems:

- **Application Layer**: Supports protocols like **HTTP** (web services), **SMTP** (email), and **FTP** (file transfer), essential for distributed applications.
- **Transport Layer**: TCP provides **reliable, ordered communication**, while UDP offers **low-latency**, faster communication. These are vital for different distributed systems use cases like video streaming (UDP) or transaction systems (TCP).
- **Internet Layer**: **IP routing** ensures data packets are transmitted across networks, supporting distributed communication across global networks.
- **Network Access Layer**: Handles physical transmission over local networks, enabling communication between distributed nodes.

OSI Model in Distributed Systems:

- **Application Layer**: Defines high-level services for distributed systems (e.g., file transfer, messaging).
- **Presentation Layer**: Manages data translation and encryption, useful for heterogeneous distributed systems.
- **Session Layer**: Supports session management between distributed nodes, crucial for **stateful** distributed applications.
- **Transport Layer**: Provides error-checking and reliable delivery, just like TCP in the TCP/IP model.
- **Network Layer**: Provides routing and addressing, similar to the Internet Layer in TCP/IP.
- **Data Link and Physical Layers**: Ensure error-free transmission, managing local network communication between distributed systems.

4. Key Differences in Supporting Distributed Systems

- **Practicality in Distributed Systems**: The **TCP/IP model** is widely used in real-world distributed systems such as **cloud computing**, **peer-to-peer networks**, and **web-based services**. Its simpler structure and practicality make it the standard for building distributed systems at scale.

- **Granularity and Flexibility**: The **OSI model** provides greater detail and flexibility with its distinct layers, which helps in conceptualizing complex aspects of distributed systems (e.g., data encryption at the Presentation layer, session management). However, its complexity makes it less common in actual deployments.

Q4.

Remote Procedure Call (RPC) in Distributed Systems

Remote Procedure Call (RPC) is a communication protocol used in distributed systems that allows a program to execute procedures (or functions) on a remote server as if they were local calls. It abstracts the complexities of network communication, making it easier for developers to design distributed applications.

Key Concepts of RPC

- **Simplification of Communication**: In a distributed system, RPC allows one program to invoke procedures or methods on another program located on a different machine without worrying about the details of network communication (e.g., data transmission, handling responses). The complexity of sending data over a network is hidden, making the process appear like a **normal function call**.

- **Client-Server Interaction**: RPC follows the **client-server model** where the client sends a request to the server to execute a procedure, and the server processes the request and sends back the result.

Key Components of RPC

1. **Client**:
   - The part of the system that makes the **RPC call**.
   - Sends a request to the server to execute a remote procedure.
   - Waits for the server to process the request and return the result.

2. **Client Stub**:

   - A **proxy** for the client that prepares (or **marshals**) the procedure call's parameters into a message that can be transmitted over the network.
   - Unpacks the server's response and passes the result to the client.

3. **RPC Runtime/Communication Module**:

   - Manages the **low-level details** of the network communication, including the sending and receiving of messages between client and server.
   - Ensures data transmission reliability, handles errors, and retransmits lost messages.

4. **Server Stub**:

   - Receives the incoming request from the client, **unmarshals** the parameters, and calls the appropriate server-side procedure.
   - After the procedure is executed, it marshals the result and sends it back to the client.

5. **Server**:

   - The server contains the actual implementation of the procedure that the client wants to execute.
   - Processes the request and returns the result to the client through the server stub.



Figure 4-5 RPC execution
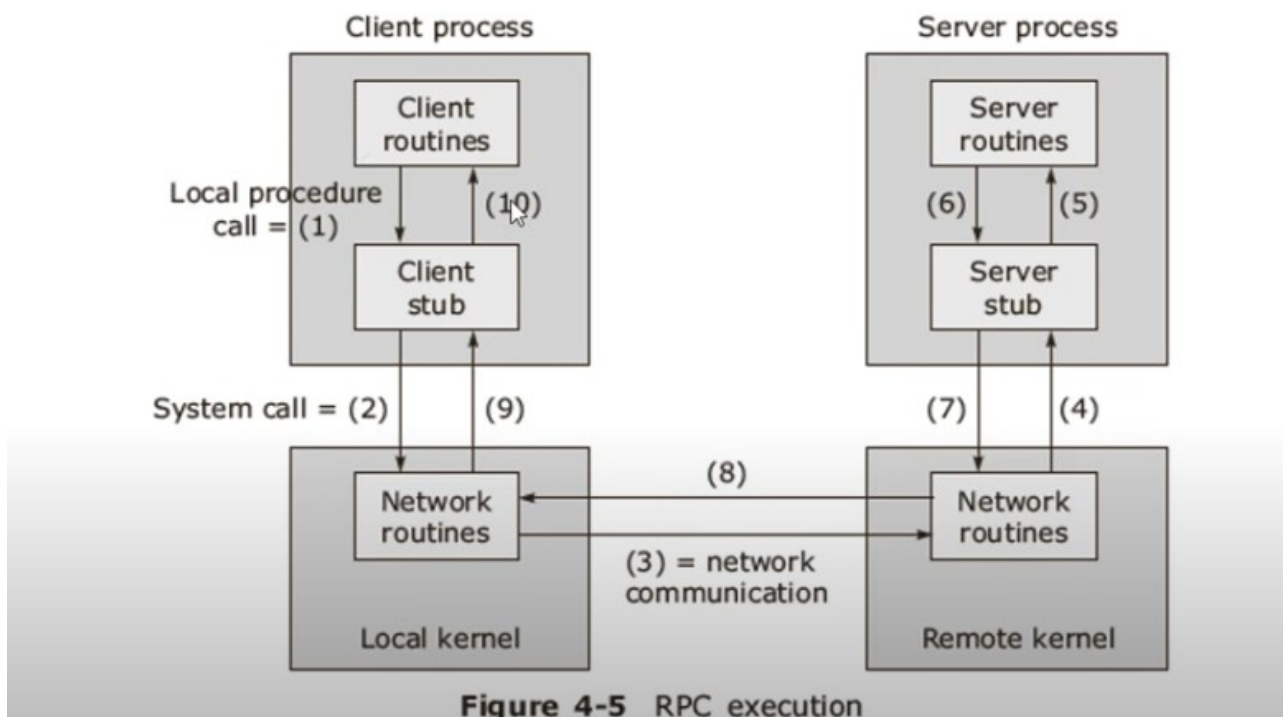
Workflow of RPC:

1. The **client** invokes a procedure using the client stub.
2. The **client stub** marshals the procedure parameters and sends them as a message to the server.

3. The **RPC runtime** on the client transmits the message over the network.
4. The **server stub** receives the message, unmarshals the parameters, and passes them to the server.
5. The **server** executes the procedure and sends the result back through the server stub.
6. The **client stub** receives the result, unmarshals it, and passes it to the client.

RPC operation

Wait for result

Client

Call remote procedure   Return from call

Request   Reply

Server

Call local procedure and return results   Time

Figure 4-4   A typical RPC

Benefits of RPC in Distributed Systems

1. **Simplified Communication**: RPC hides the complexities of network programming, such as sockets and data transmission protocols, allowing developers to focus on business logic instead of networking details.

2. **Transparency**: The client code looks almost identical to a local function call, making distributed systems easier to build and maintain.

3. **Interoperability**: RPC can be implemented across different systems and languages, allowing programs written in different languages or running on different platforms to communicate effectively.

4. **Scalability**: RPC allows the server-side computation to be distributed across multiple machines, enhancing system scalability.

Q5.

Group Communication in Distributed Systems

Group communication plays a critical role in **distributed systems**, where multiple processes or nodes need to communicate and collaborate. It enables **communication among multiple participants simultaneously**, which is essential for tasks such as **replication**, **fault tolerance**, **synchronization**, and **broadcasting messages** across nodes. Group communication allows nodes in a distributed environment to **coordinate actions, share data, and maintain consistency**.

In this context, group communication ensures that:

- **Messages are delivered to multiple recipients simultaneously** (multicasting or broadcasting).
- **Order of message delivery** is maintained, ensuring **consistency** across distributed nodes.
- It helps in **load balancing**, **failure detection**, and **system recovery** by allowing nodes to share status and data.

Significance of Group Communication:

1. **Fault Tolerance**: Group communication ensures that even if some nodes fail, the message can still reach other nodes in the system.
2. **Consistency**: It maintains a consistent view across distributed nodes, crucial for replication and distributed databases.
3. **Scalability**: Efficient group communication protocols ensure that the system can scale to handle a large number of nodes while still maintaining communication efficiency.
4. **Coordination**: Many distributed algorithms rely on group communication to coordinate actions between nodes, ensuring synchronized operations.
5. **System Recovery**: In case of failure, group communication allows surviving nodes to detect and take over tasks from failed nodes.
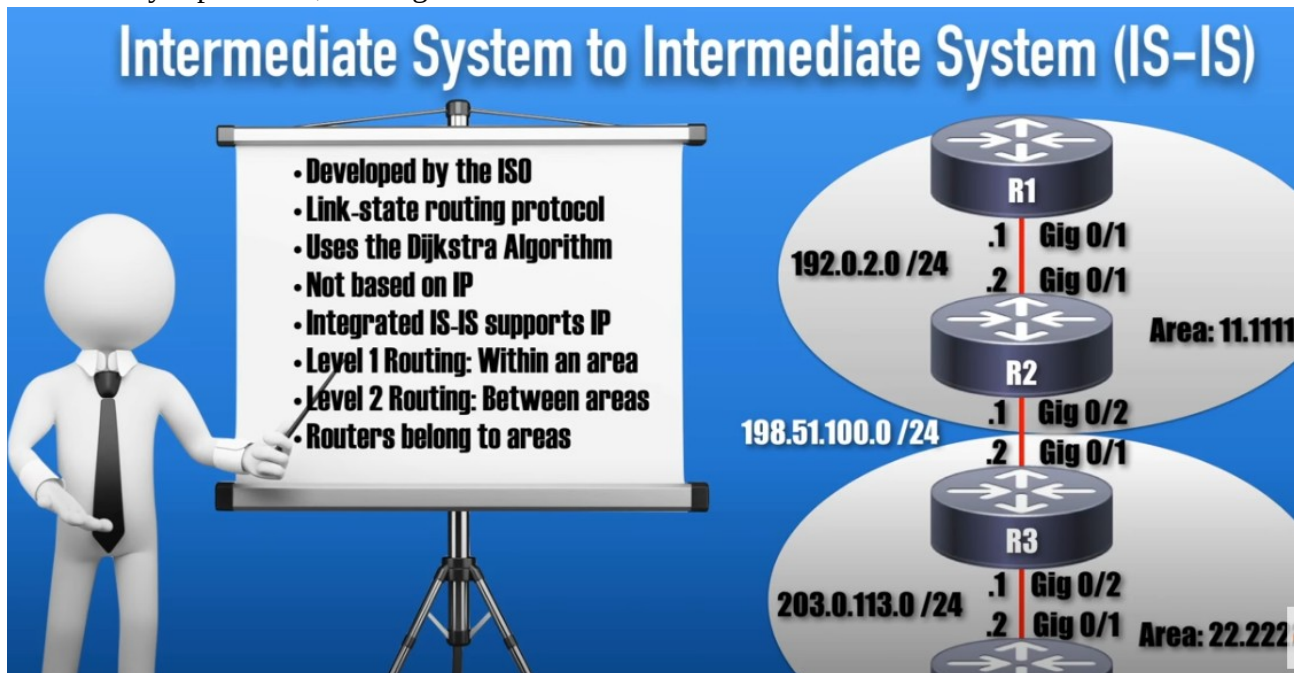

IS-IS Protocol and its Role in Group Communication

The **Intermediate System to Intermediate System (IS-IS)** protocol is a **link-state routing protocol** used primarily to move information efficiently across networks. Originally designed for **OSPF (Open Shortest Path First)**, IS-IS has become essential in managing **large-scale distributed systems**, particularly in telecommunications and data centers.

Key Features of IS-IS Protocol:

1. **Link-State Routing**: IS-IS uses **link-state information** to create a network topology map and determine the shortest path for data transmission.
2. **Scalability**: IS-IS is designed to scale efficiently in large networks with many routers and nodes, making it well-suited for **distributed systems**.
3. **Hierarchical Design**: IS-IS supports **multi-level hierarchies**, allowing networks to be organized into regions for more efficient routing, reducing overhead in group communication.

4. **Multiprotocol Support**: IS-IS is not limited to just IP but can also handle other network-layer protocols, making it flexible for various distributed architectures.

IS-IS in Group Communication:

The IS-IS protocol enhances group communication in distributed systems by providing the following capabilities:

1. **Efficient Multicast Communication**: IS-IS supports the transmission of **multicast** traffic, enabling the same message to be sent to multiple destinations simultaneously, which is crucial for **group communication**.
2. **Reliability through Link-State Database**: By maintaining a detailed **link-state database**, IS-IS ensures that all nodes have a consistent view of the network topology. This is important for group communication, where consistency across nodes is essential.
3. **Fast Convergence**: In the event of network changes, IS-IS can quickly converge to update routing information, ensuring that messages in group communication are always delivered along the most efficient path.
4. **Failure Recovery**: In case of node or link failures, IS-IS reroutes traffic through alternate paths, ensuring that messages still reach all participants in the group.

IS-IS Protocol Operation:

1. **Hello Packets**: Routers send **Hello packets** to discover neighbors and establish adjacencies.
2. **Link-State Packets (LSPs)**: Once adjacencies are formed, routers exchange **LSPs**, which contain information about their links. These LSPs are propagated throughout the network to build a complete topology map.
3. **SPF Algorithm**: IS-IS uses the **Shortest Path First (SPF)** algorithm to calculate the best paths to each destination, ensuring that data is sent through the most efficient routes.
4. **Multicast Traffic Management**: IS-IS ensures that multicast traffic reaches all intended recipients in a group without redundancy or loss, maintaining the efficiency of group communication.

Q6.

# ATM (Asynchronous Transfer Mode) and its Relevance to Distributed Systems

Asynchronous Transfer Mode (ATM) is a high-speed, connection-oriented networking technology designed to efficiently transmit a variety of data types (voice, video, and data) over a single network. It uses fixed-sized cells (53 bytes), which makes it ideal for real-time data transmission, ensuring consistent and predictable performance.

In the context of **distributed systems**, ATM plays a significant role in providing reliable and efficient communication between distributed nodes, as it offers features that support high-speed and scalable communication necessary for distributed applications.

Working of ATM:
ATM standard uses two types of connections. i.e., Virtual path connections (VPCs) which consist of Virtual channel connections (VCCs) bundled together which is a basic unit carrying a single stream of cells from user to user. A virtual path can be created end-to-end across an ATM network, as it does not rout the cells to a particular virtual circuit. In case of major failure, all cells belonging to a particular virtual path are routed the same way through the ATM network, thus helping in faster recovery.

Switches connected to subscribers use both VPIs and VCIs to switch the cells which are Virtual Path and Virtual Connection switches that can have different virtual channel connections between them, serving the purpose of creating a virtual trunk between the switches which can be handled as a single entity. Its basic operation is straightforward by looking up the connection value in the local translation table determining the outgoing port of the connection and the new VPI/VCI value of connection on that link.
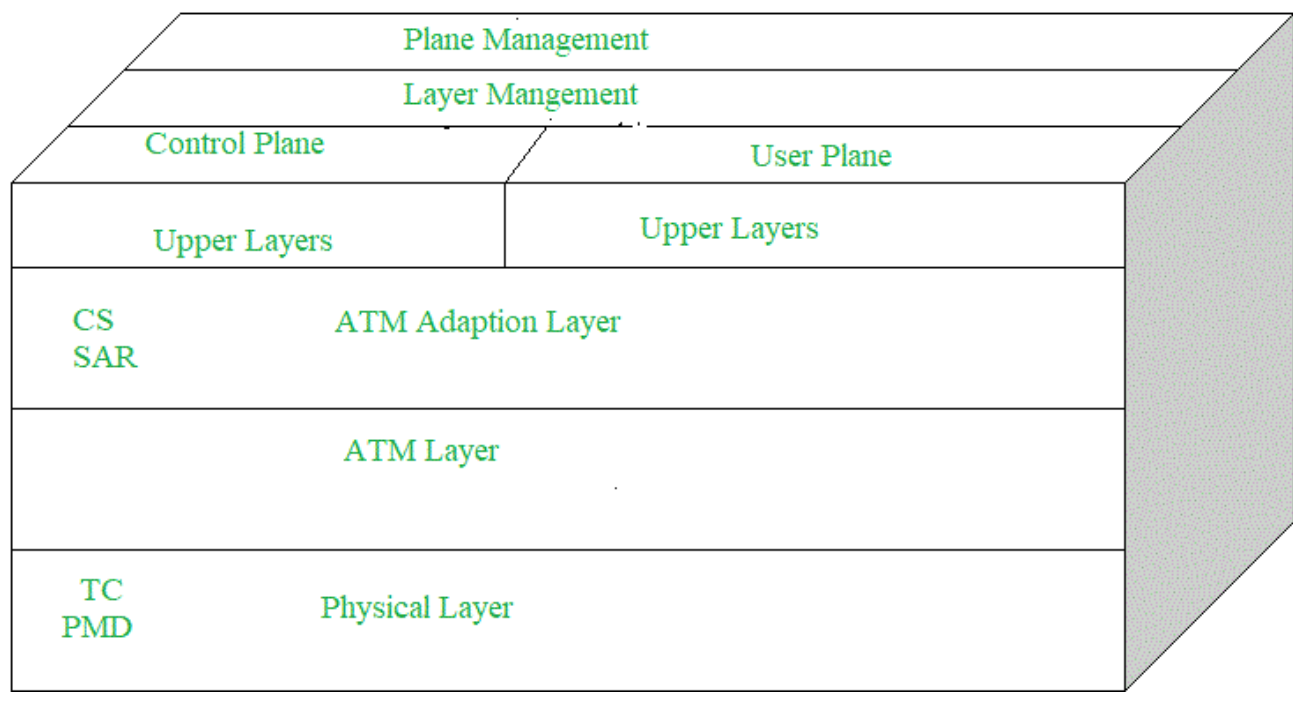
Advantages of ATM in Distributed Systems

1. **High Performance and Low Latency**:

- **Fixed-size cells (53 bytes)**: ATM transmits data in small, fixed-sized cells, which enables quick processing at switches and routers, reducing delay and ensuring **low latency**.
- **Efficient handling of real-time traffic**: Distributed systems often deal with diverse traffic types like voice, video, and data. ATM's ability to support **real-time data transmission** ensures that such traffic can be handled efficiently, making it suitable for applications like **multimedia streaming**, **video conferencing**, and **real-time data analytics**.

2. **Scalability and High Throughput**:

- ATM supports **very high data rates** (up to Gbps), which is crucial for distributed systems operating over large networks or dealing with a high volume of data exchange. Distributed systems with heavy data traffic or many interconnected nodes can scale better using ATM due to its high **bandwidth capacity**.
- **Scalability** is achieved because ATM is well-suited for a variety of network environments, from **local area networks (LANs)** to **wide area networks (WANs)**, ensuring that it can handle growing communication needs of distributed systems.

3. **Quality of Service (QoS)**:

- ATM offers **Quality of Service (QoS)** guarantees, enabling distributed systems to define priorities for different types of data. For example, real-time applications can request higher priority than non-critical data, ensuring timely delivery of essential information in the system.
- This QoS feature is particularly important for **time-sensitive applications** in distributed systems, such as **real-time financial transactions** or **remote medical procedures**.

4. **Reliable Data Transmission**:

- **Connection-Oriented Communication**: ATM establishes a virtual circuit between nodes before data transmission, ensuring **reliable delivery** of data. This is especially valuable in distributed systems where **fault tolerance** and **consistent communication** are crucial.
- The **virtual circuit** mechanism ensures that packets follow the same route, reducing the chances of data arriving out of order, which is critical for ensuring the **reliability** of distributed applications.

5. **Adaptability for Various Data Types**:

- ATM supports **integrated data transfer**, meaning it can handle different data types (voice, video, and regular data) within the same network infrastructure. Distributed systems often need to handle multiple types of traffic simultaneously, and ATM's ability to **prioritize and integrate different data types** makes it highly adaptable.
- This makes ATM suitable for applications like **distributed multimedia systems**, **scientific data distribution**, and **remote collaboration tools**.

Q7.

Motivation for Distributed Systems Development:

1. **Scalability:** Traditional centralized systems face limitations in scaling to handle increasing loads. Distributed systems enable horizontal scaling by adding more nodes to a network, efficiently managing larger data volumes and processing demands.

2. **Fault Tolerance:** A single point of failure in centralized systems can cause a complete system crash. Distributed systems improve reliability by distributing tasks across multiple machines, ensuring that if one node fails, others continue functioning.

3. **Resource Sharing:** Distributed systems allow for the sharing of computational resources, storage, and data across different geographic locations, maximizing resource utilization and reducing costs.

4. **Performance and Speed:** By distributing tasks among multiple machines, distributed systems can process large-scale computations in parallel, reducing latency and improving overall performance.

5. **Geographical Distribution:** In global applications, data and services need to be accessible across different regions. Distributed systems ensure that users worldwide can access data and services with minimal delay.

Problems Addressed:

- **Scalability and resource bottlenecks in centralized systems.**
- **Single points of failure causing system outages.**
- **Slow performance due to large computational tasks.**
- **Challenges in global access to resources.**

**Compare Centralized n/w here too

Q8.

Goals :

1.**Making Resources Accessible**

2.**Distribution transparency :** hide the fact that its processes and resources are physically distributed across multiple computers.

Types of Transparency

Transparency	Description
Access	Hide differences in data representation and how a resource is accessed
Location	Hide where a resource is located
Migration	Hide that a resource may move to another location
Relocation	Hide that a resource may be moved to another location while in use
Replication	Hide that a resource is replicated
Concurrency	Hide that a resource may be shared by several competitive users
Failure	Hide the failure and recovery of a resource

Figure 1-2. Different forms of transparency in a distributed system (ISO, 1995).

Degree of Transparency

3.**Openness** : An open distributed system is a system that offers services according to standard rules that describe the syntax and semantics of those services

4.**Scalability** : Size,geographic location , administratively scalable , Scaling Techniques

5.**Pitfalls**:

> 1. The network is reliable.
>
> 2. The network is secure.
>
> 3. The network is homogeneous.
>
> 4. The topology does not change.
>
> 5. Latency is zero.
>
> 6. Bandwidth is infinite.
>
> 7. Transport cost is zero.

8. There is one administrator.\

Advantages over Centralized Systems:

1. **Improved Reliability and Availability:** Distributed systems reduce the risk of total system failure by distributing tasks across multiple nodes. If one node fails, others can take over, ensuring the system remains operational.

2. **Enhanced Scalability:** Unlike centralized systems, distributed systems can scale horizontally by adding more machines, enabling them to handle increasing loads without performance degradation.

3. **Better Resource Sharing:** Distributed systems allow for the sharing of resources (e.g., processing power, storage) across multiple locations, improving efficiency and optimizing resource utilization globally.

Q9.

Disadvantages or Challenges of Distributed Systems:

1. **Complexity in Design and Maintenance:** Distributed systems involve coordinating multiple machines, which adds complexity to both their design and ongoing maintenance. Managing data consistency, communication, and synchronization between nodes requires sophisticated algorithms and protocols, making the system more challenging to develop and maintain.

2. **Security Risks:** With data and processes distributed across various locations, security becomes a significant concern. Distributed systems are more vulnerable to attacks like unauthorized access, data breaches, and denial of service (DoS), requiring robust security measures like encryption, secure communication protocols, and access control mechanisms.

3. **Network Reliability and Latency Issues:** Distributed systems rely heavily on network connectivity. Any delays or interruptions in the network can degrade performance or cause failures in the system. Additionally, achieving low-latency communication between geographically distributed nodes can be difficult, impacting system efficiency.

Impact on Adoption:

- **Higher Costs and Expertise Requirements:** The complexity and need for specialized expertise can make distributed systems expensive to develop, implement, and maintain, slowing down adoption for smaller organizations.
- **Security Concerns:** The heightened risk of security vulnerabilities can deter industries handling sensitive data from adopting distributed systems without substantial security investment.
- **Dependence on Reliable Networks:** In areas with unstable or limited network infrastructure, adopting distributed systems can be challenging, as network issues can compromise their effectiveness.

Q11

Key Software Concepts in Distributed Systems:

1. **Operating Systems (OS):** The OS in distributed systems manages the hardware resources (CPU, memory, storage) of individual machines. It ensures that processes on each node can execute efficiently and handle tasks like process scheduling, memory management, and I/O operations. In a distributed environment, the OS also handles communication between nodes, supporting network protocols and security mechanisms to ensure reliable communication across the system.

2. **Middleware:** Middleware is the software layer that sits between the operating system and application software in distributed systems. It abstracts the complexity of the underlying network, making it easier for applications to interact with different parts of the system. Middleware provides essential services like message passing, remote procedure calls (RPCs), and data synchronization, ensuring smooth communication, coordination, and resource sharing among distributed components.

3. **Application Software:** This refers to the end-user applications or services running on top of the distributed system. These applications leverage the distributed nature of the system to offer scalable, reliable, and fault-tolerant services, such as web applications, cloud services, or distributed databases. The application software interacts with both the OS and middleware to execute tasks and provide seamless experiences to users.

Role of These Components in Distributed Systems:

- **Operating System:** It ensures the proper execution of processes on each node, handling hardware-level tasks and communication across the distributed network.

- **Middleware:** It facilitates coordination and communication between distributed components, abstracting complexity and providing a unified interface for applications.

- **Application Software:** It delivers distributed services to users by leveraging the capabilities provided by the OS and middleware to perform tasks like data storage, processing, and interaction across multiple nodes.