

# Distributed Systems

Introduction



# Slides Credits

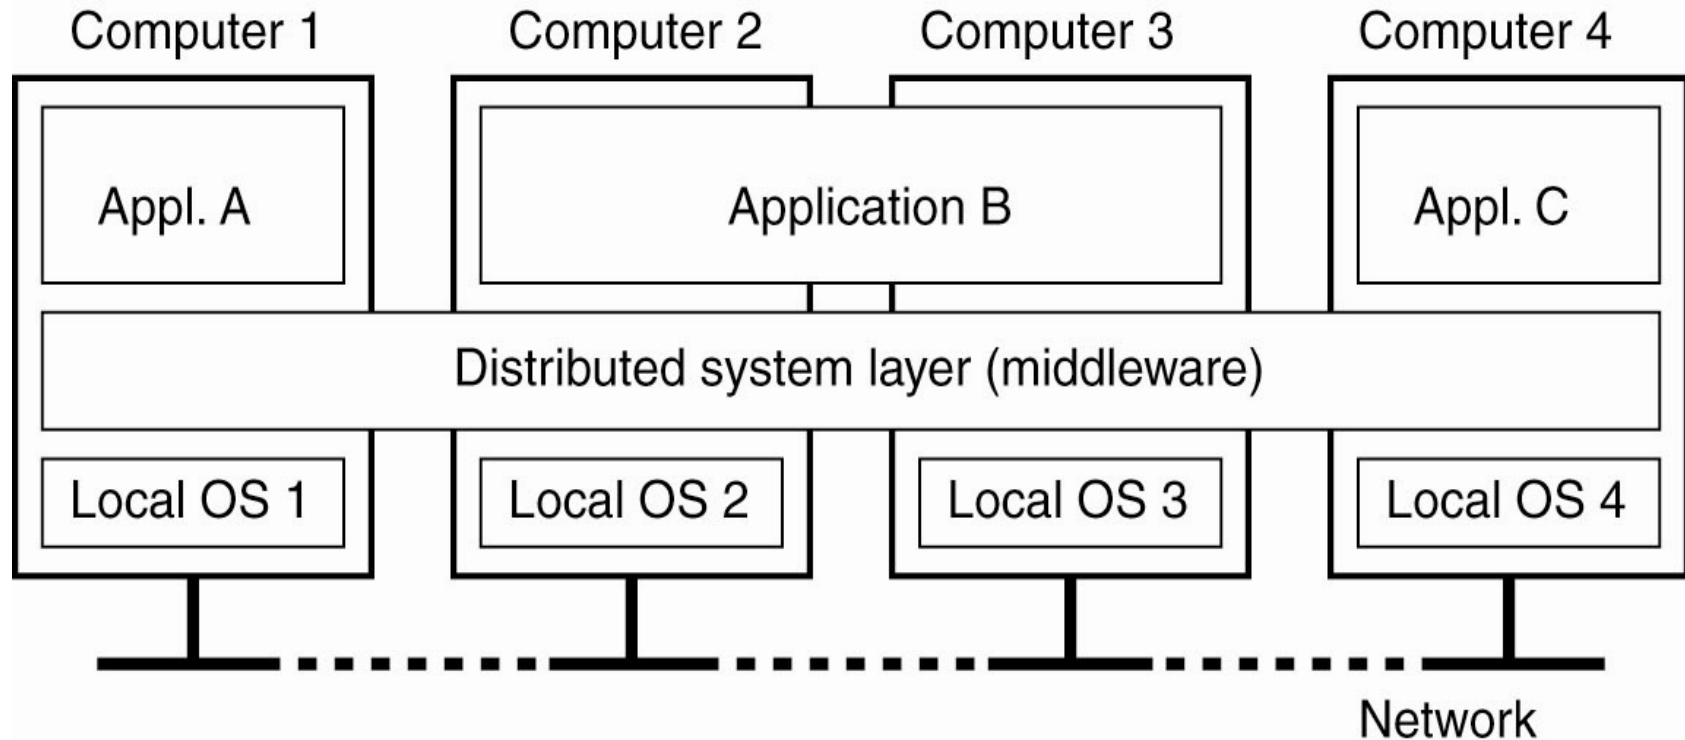
**Note:** Most of the course presentations are based on those developed by Andrew S. Tanenbaum and Maarten van Steen.

# Definition of a Distributed System

A distributed system is:

A collection of independent computers  
that appears to its users  
as a single coherent system

# Definition of a Distributed System



A distributed system organized as middleware. The middleware layer extends over multiple machines, and offers each application the same interface

# Goals of Distributed Systems

- Easily Connect Users/Resources
- Exhibit Distribution Transparency
- Support Openness
- Be Scalable:
  - in size
  - geographically
  - administratively

Looking at these goals helps us answer the question: “Is building a distributed system worth the effort?”

# Transparency in a Distributed System

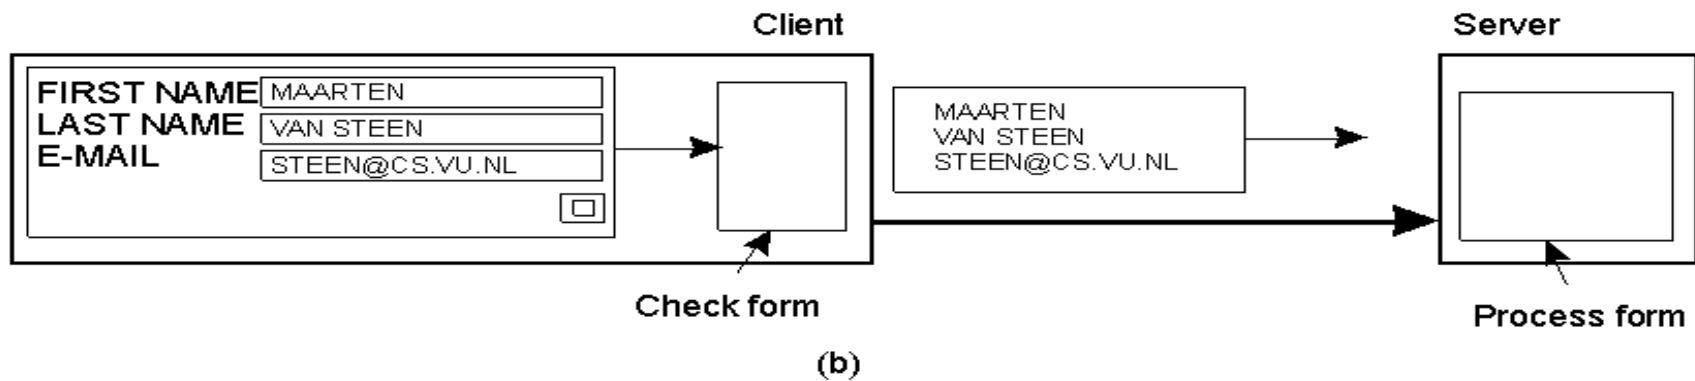

Transparency	Description
Access	Hide differences in data representation and how a resource is accessed
Location	Hide where a resource is located
Migration	Hide that a resource may move to another location
Relocation	Hide that a resource may be moved to another location while in use
Replication	Hide that a resource is replicated
Concurrency	Hide that a resource may be shared by several competitive users
Failure	Hide the failure and recovery of a resource

# Scalability Limitations

Concept	Example
Centralized services	A single server for all users
Centralized data	A single on-line telephone book
Centralized algorithms	Doing routing based on complete information

Examples of scalability limitations

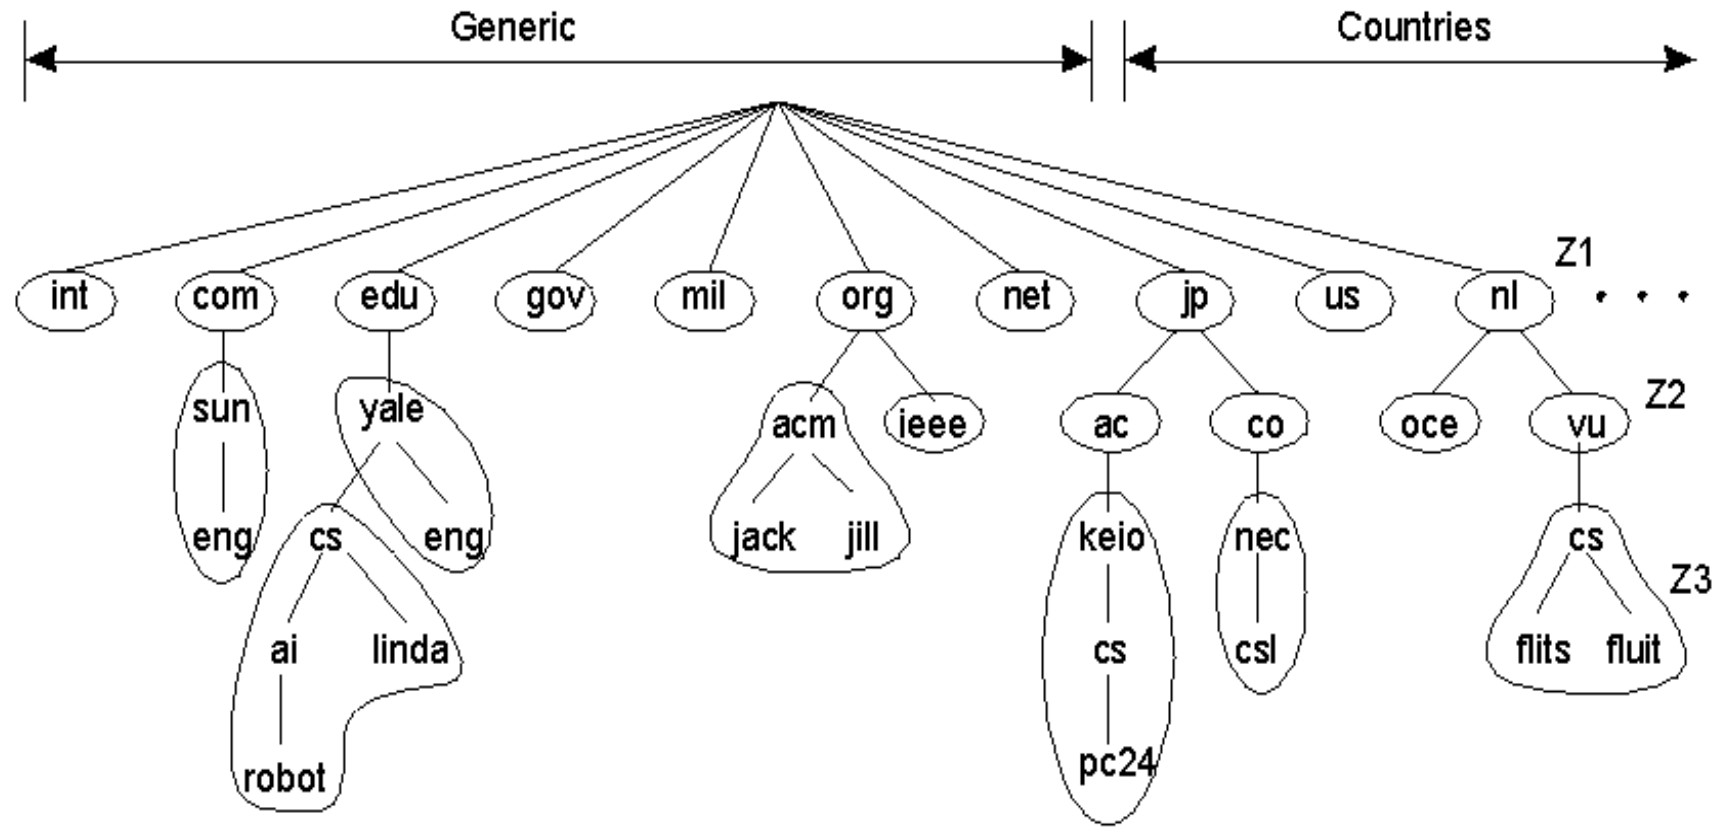
# Scaling Techniques (1)



The difference between letting (a) a server or (b) a client check forms as they are being filled



# Scaling Techniques (2)



An example of dividing the DNS name space into zones

## :Characteristics of decentralized algorithms

- No machine has complete information about the system state.
- Machines make decisions based only on local information.
- Failure of one machine does not ruin the algorithm.
- There is no implicit assumption that a global clock exists.


# Pitfalls when Developing Distributed Systems

- The network is reliable
- The network is secure
- The network is homogeneous
- The topology does not change
- Latency is zero
- Bandwidth is infinite
- Transport cost is zero
- There is one administrator

# Types of Distributed Systems

- Distributed Computing Systems
  - High Performance Computing (HPC)
- Distributed Information Systems
  - Transaction Processing Systems (TPS)
  - Enterprise Application Integration (EAI)
- Distributed Pervasive Systems
  - Ubiquitous Systems

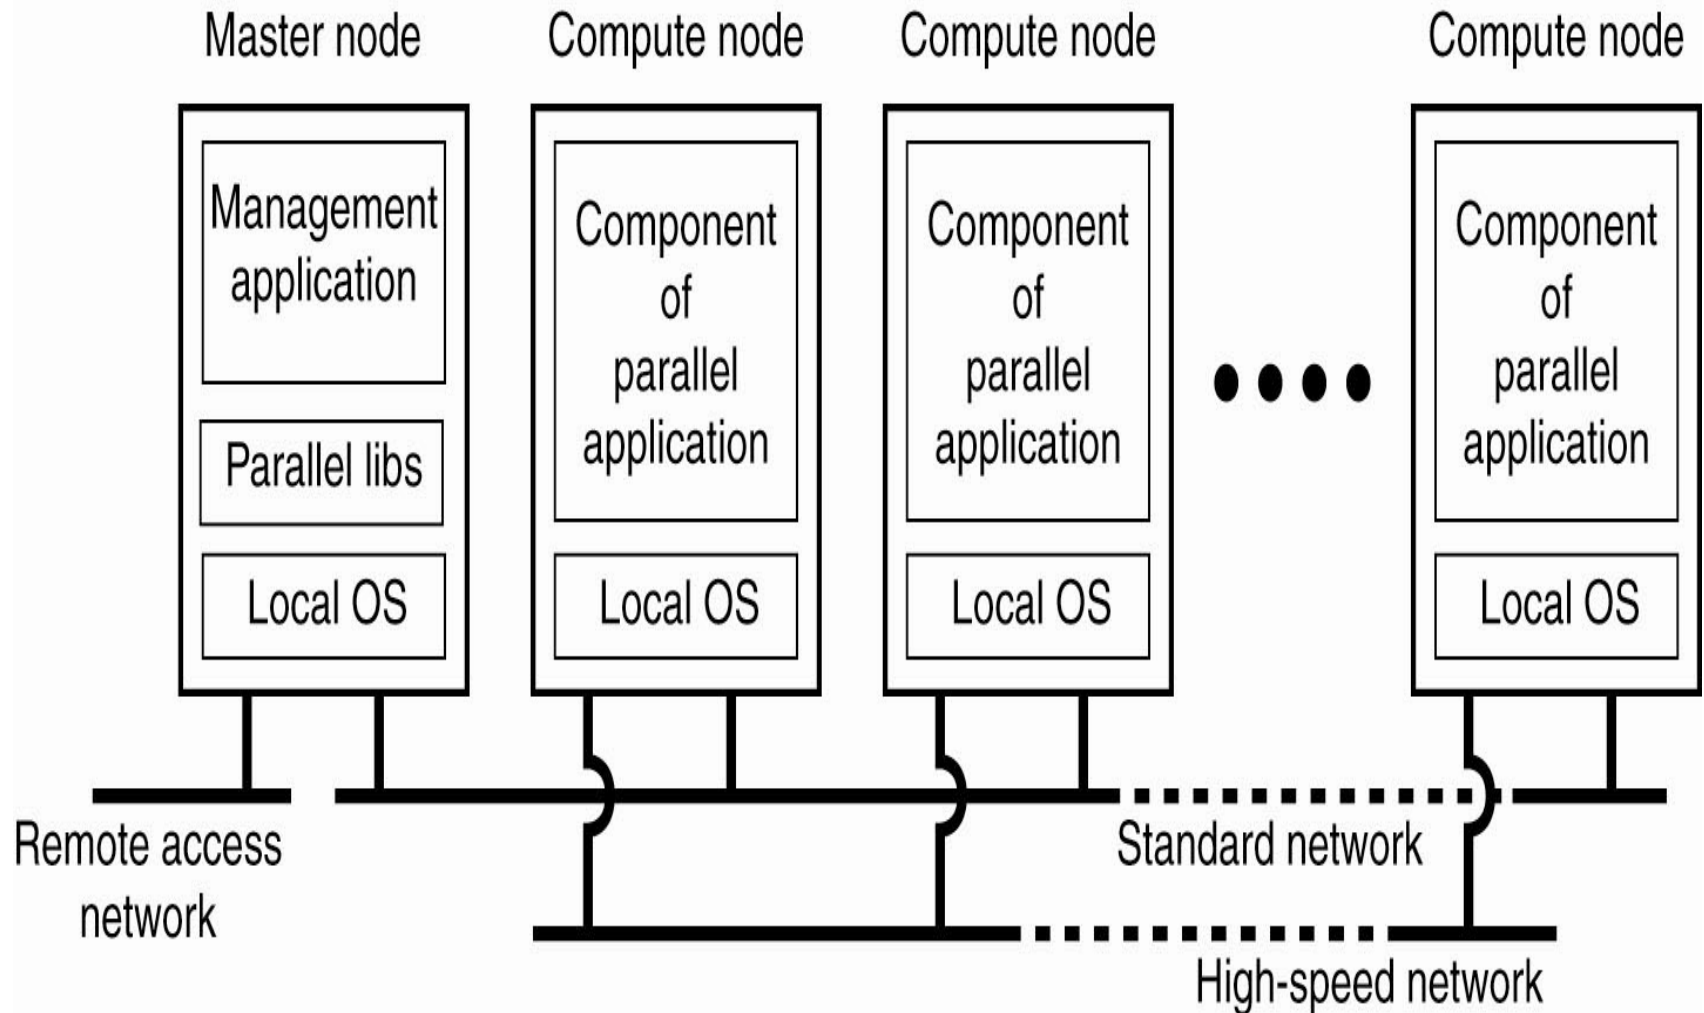
# Clustered Systems Architecture



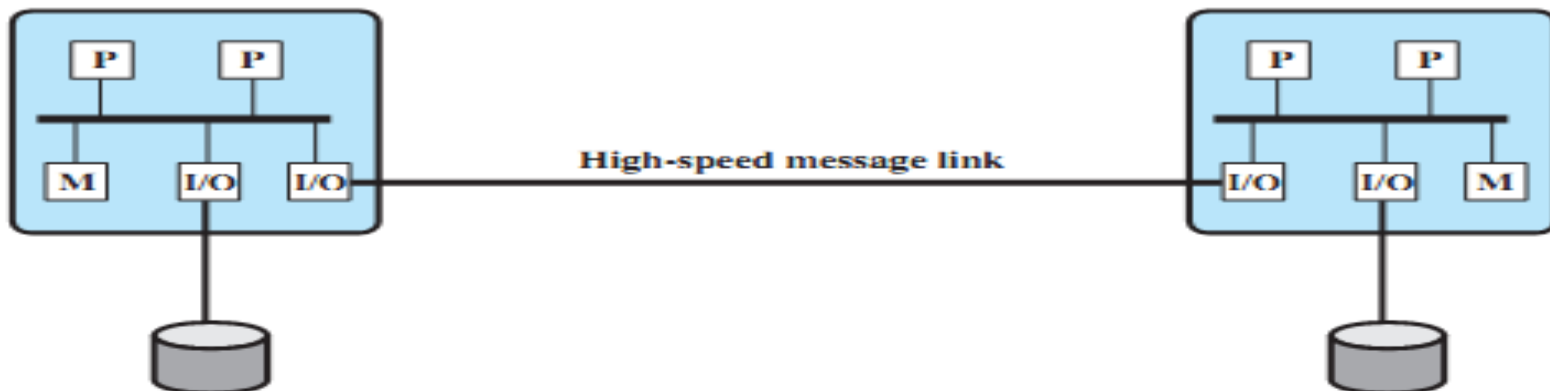
# Cluster Computing Systems

- Collection of similar workstations/PCs, closely connected by means of a high-speed LAN:
  - Each node runs the same OS.
  - Homogeneous environment
  - Can serve as a supercomputer
  - Excellent for parallel programming
- Examples: Linux-based Beowulf clusters, MOSIX (from Hebrew University).


# Architecture for Cluster Computing System



# Cluster Configurations



(a) Standby server with no shared disk






# Grid Computing Systems

- Collection of computer resources, usually owned by multiple parties and in multiple locations, connected together such that users can share access to their combined power:
  - Can easily span a wide-area network
  - Heterogeneous environment
  - Crosses administrative/geographic boundaries
  - Supports Virtual Organizations (VOs)
- Examples: EGEE - Enabling Grids for E-Science (Europe), Open Science Grid (USA).

# Architecture for Grid Computing Systems



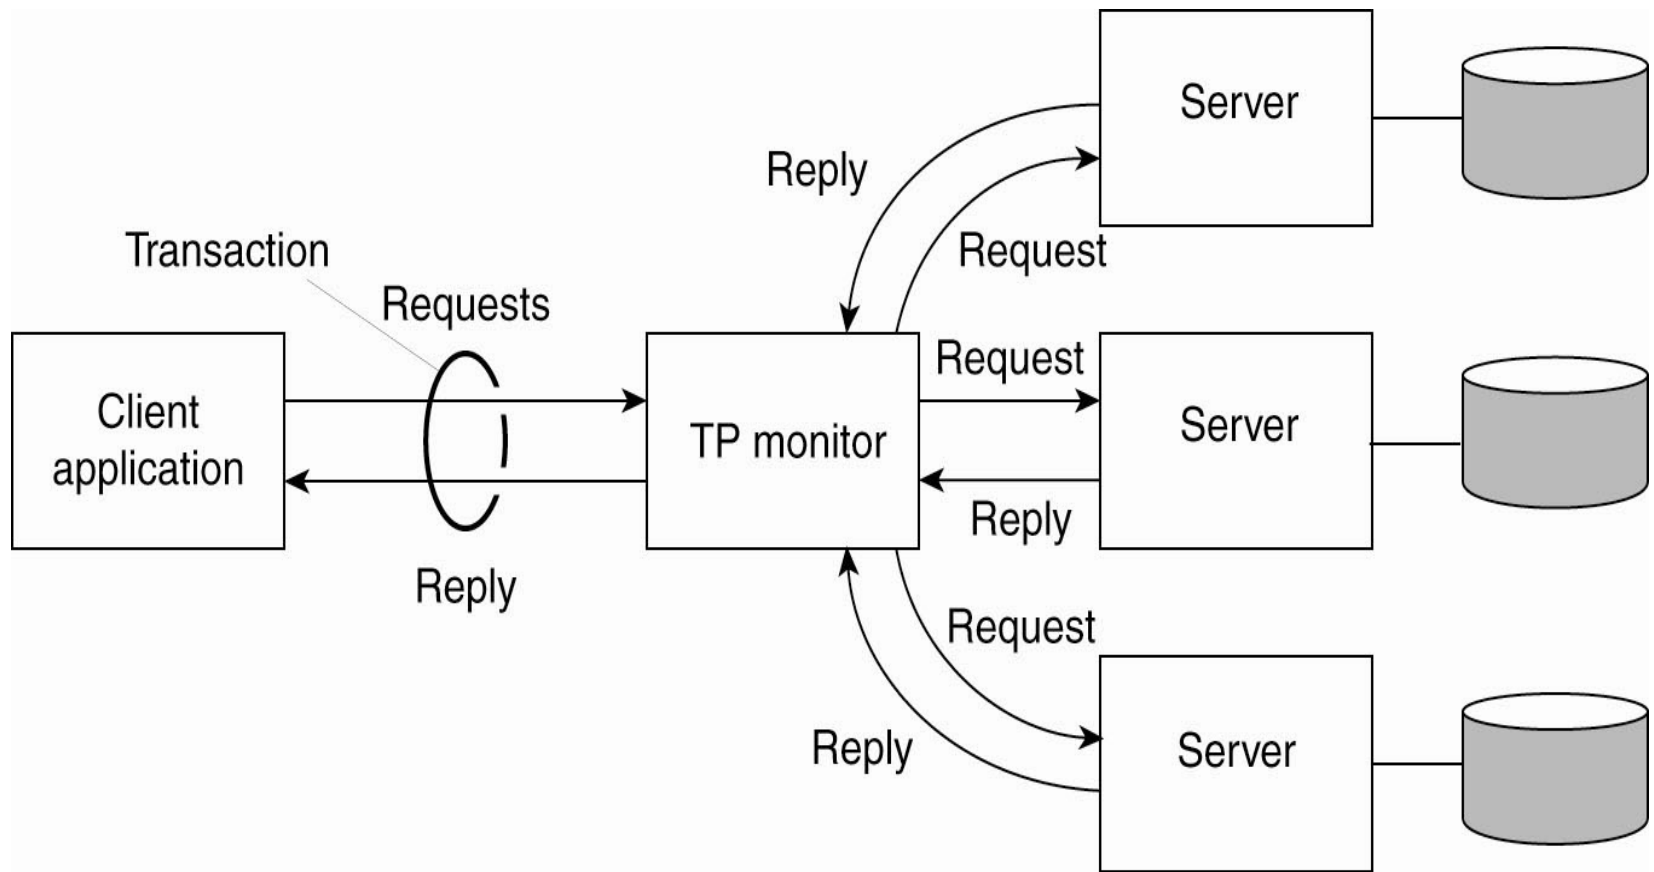
# Cloud Computing Systems (1)

- Collection of computer resources, usually owned by a single entity, connected together such that users can lease access to a share of their combined power:
  - **Location independence**: the user can access the desired service from anywhere in the world, using any device with any (supported) system.
  - **Cost-effectiveness**: the whole infrastructure is owned by the provider and requires no capital outlay by the user.
  - **Reliability**: enhanced by way of multiple redundant sites, though outages can occur, leaving users unable to remedy the situation.

## Cloud Computing Systems (2)


- **Scalability**: user needs can be tailored to available resources as demand dictates – cost benefit is obvious.
- **Security**: low risk of data loss thanks to centralization, though problems with control over sensitive data need to be solved.
- **Readily consumable**: the user usually does not need to do much deployment or customization, as the provided services are easy to adopt and ready-to-use.
- Examples: Amazon EC2 (Elastic Compute Cloud), Google App Engine, IBM Enterprise Data Center, MS Windows Azure, SUN Cloud Computing.

# Transaction Processing Systems (TPS)



The role of a TP monitor in distributed systems

# Enterprise Application Integration



# Communication Middleware Models/Paradigm

- Distributed File Systems
- Remote Procedure Call (RPC)
- Distributed Objects (RMI)
- Distributed Documents

# Distributed Pervasive Systems


- Requirements for pervasive systems:
  - Embrace contextual changes
  - Encourage ad hoc composition
  - Recognize sharing as the default
  - Support distribution transparency



# Electronic Health Care Systems (1)

- Questions to be addressed for health care systems:
  - Where and how should monitored data be stored?
  - How can we prevent loss of crucial data?
  - What infrastructure is needed to generate and propagate alerts?
  - How can physicians provide online feedback?
  - How can extreme robustness of the monitoring system be realized?
  - What are the security issues and how can the proper policies be enforced?

# Electronic Health Care Systems (2)

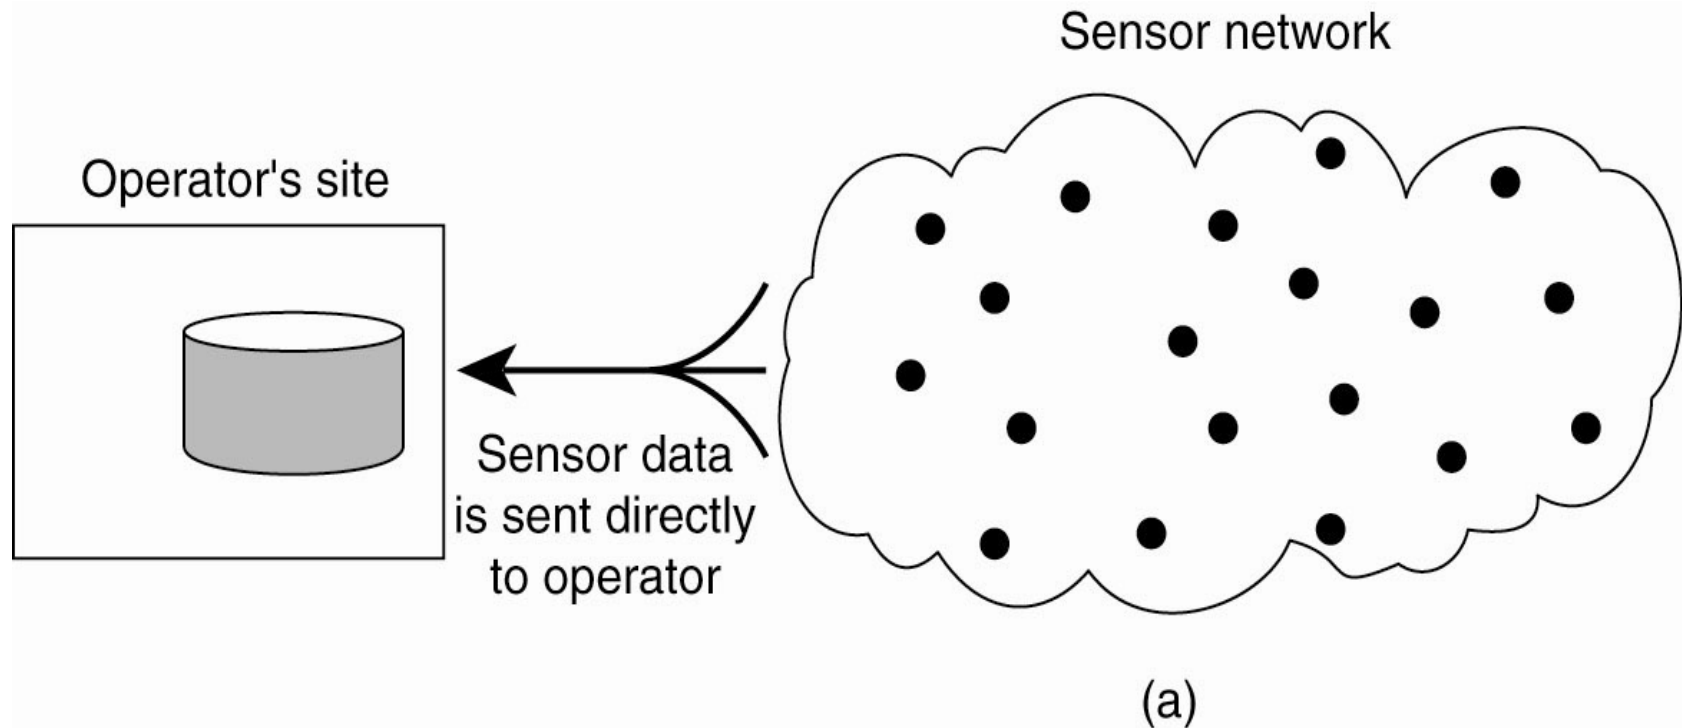


Monitoring a person in a pervasive electronic health care system, using (a) a local hub or (b) a continuous wireless connection

# Sensor Networks (1)

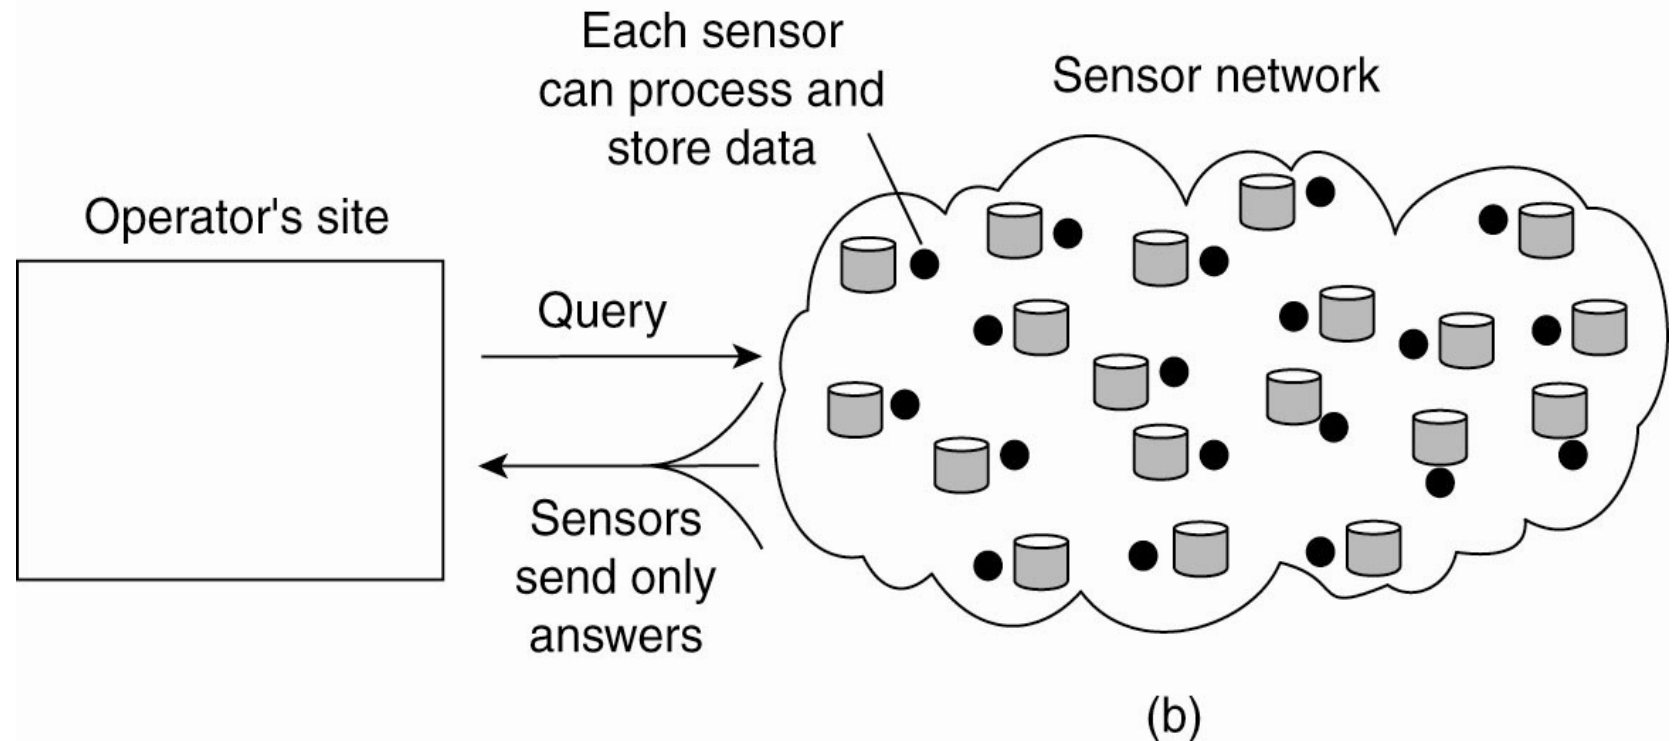
- The nodes to which sensors are attached are:
  - Many (10s-1000s).
  - Simple (i.e., hardly any memory, CPU power, or communication facilities).
  - Often battery-powered (or even battery-less).
- Questions concerning sensor networks:
  - How do we (dynamically) set up an efficient tree in a sensor network?
  - How does aggregation of results take place?  
Can it be controlled?
  - What happens when network links fail?

## Sensor Networks (2)



Organizing a sensor network database, while storing and processing data (a) only at the operator's site or ...

## Sensor Networks (3)



Organizing a sensor network database, while storing and processing data ... or (b) only at the sensors